

# Week 02, Monday

## Memory and Data

### The C View of Data

2/33

A C program sees data as a collection of *variables*

```
int g = 2;

int main(void)
{
    int i;
    int v[5]
    char *s = "Hello";
    int *n = malloc(sizeof(int));
    ...
}
```

Each variable has a number of properties (e.g. name, type, size)

### ... The C View of Data

3/33

Variables are examples of *computational objects*

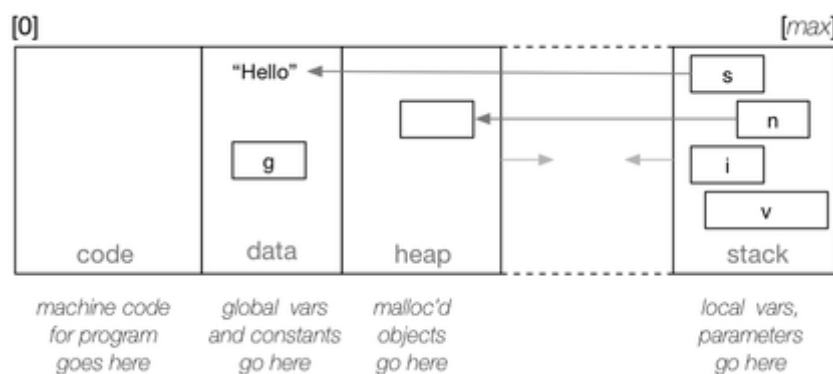
Each computational object has

- a *location* in memory
- a *value* (ultimately just a bit-string)
- a *name* (unless created by `malloc()`)
- a *type*, which determines ...
  - its *size* (in units of whole bytes, `sizeof`)
  - how to *interpret* its value
  - what *operations* apply to the value
- a *scope* (where it's visible within the program)
- a *lifetime* (during which part of program execution it exists)

### ... The C View of Data

4/33

C allocates data objects to various well-defined regions of memory during program execution



### Exercise 1: Properties of Variables

5/33

Identify the properties of each of the named objects in the following:

```
int a;           // global int variable

int main(void) {
    int b;       // local int variable
    char c;      // local char variable
    char d[10];  // local char array
    ...
}


int e;           // global? int variable

int f(int g) {   // function + parameter
    double h;    // local double variable
    ...
}
```

## The Physical View of Data

6/33

Memory = indexed array of bytes



Indexes are "memory addresses" (a.k.a. pointers)

Data can be fetched in chunks of 1,2,4,8 bytes

## Memory

7/33

Also called: RAM, main memory, primary storage, ...

Technology: semiconductor-based

Distinguishing features

- relatively large (e.g.  $2^{28}$  bytes)
- any byte can be fetched with same cost
- cost of fetching 1,2,4,8 bytes is small (ns)

Two properties related to data persistence

- *volatile* (e.g. DRAM) ... data lost when powered off
- *non-volatile* (e.g. EEPROM) ... data stays when powered off

### ... Memory


8/33

When addressing objects in memory

- any byte address can be used to fetch 1-byte object
- byte address for  $N$ -byte object must be divisible by  $N$

Data is fetched into  $N$ -byte CPU registers for use

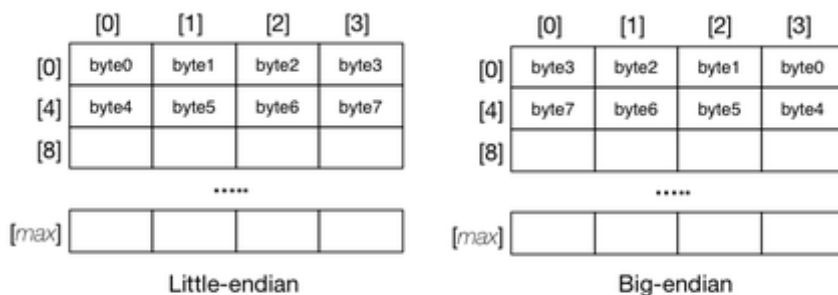
Data bytes in registers may be in different order to memory, e.g.



## ... Memory

9/33

Memories can be categorised as *big-endian* or *little-endian*



## Data Representation

### Data Representation

11/33

Ultimately, memory allows you to

- load bit-strings of sizes 1,2,4,8 bytes
- from  $N$ -byte boundary addresses
- into registers in the CPU

What you are presented with is a string of 8,16,32,64 bits

Need to *interpret* this bit-string as a meaningful value

*Data representations* provide a way of assigning meaning to bit-strings

### Character Data

12/33

Character data has several possible representations (encodings)

The two most common:

- ASCII (ISO 646)
  - 7-bit values, using lower 7-bits of a byte (top bit always zero)
  - can encode roman alphabet, digits, punctuation, control chars
- UTF-8 (Unicode)
  - 8-bit values, with ability to extend to multi-byte values
  - can encode all human languages plus other symbols

(e.g.  $\sqrt{\sum \forall \exists}$  or 🍌 🍌 🍌 🍌 🍌 🍌 )

## ASCII Character Encoding

13/33

Uses values in the range 0x00 to 0x7F (0..127)

Characters partitioned into sequential groups

- control characters (0..31) ... e.g. '\0', '\n'
- punctuation chars (32..47,91..96,123..126)
- digits (48..57) ... '0'..'9'
- upper case alphabetic (65..90) ... 'A'..'Z'
- lower case alphabetic (97..122) ... 'a'..'z'

In C, can map between char and ascii code by e.g. ((int)'a')

Sequential nature of groups allow for e.g. (ch - '0')

### ... ASCII Character Encoding

14/33

Hexademical ASCII char table (from `man 7 ascii`)

00 nul	01 soh	02 stx	03 etx	04 eot	05 enq	06 ack	07 bel
08 bs	09 ht	0a nl	0b vt	0c np	0d cr	0e so	0f si
10 dle	11 dc1	12 dc2	13 dc3	14 dc4	15 nak	16 syn	17 etb
18 can	19 em	1a sub	1b esc	1c fs	1d gs	1e rs	1f us
20 sp	21 !	22 "	23 #	24 \$	25 %	26 &	27 '
28 (	29 )	2a *	2b +	2c ,	2d -	2e .	2f /
30 0	31 1	32 2	33 3	34 4	35 5	36 6	37 7
38 8	39 9	3a :	3b ;	3c <	3d =	3e >	3f ?
40 @	41 A	42 B	43 C	44 D	45 E	46 F	47 G
48 H	49 I	4a J	4b K	4c L	4d M	4e N	4f O
50 P	51 Q	52 R	53 S	54 T	55 U	56 V	57 W
58 X	59 Y	5a Z	5b [	5c \	5d ]	5e ^	5f _
60 `	61 a	62 b	63 c	64 d	65 e	66 f	67 g
68 h	69 i	6a j	6b k	6c l	6d m	6e n	6f o
70 p	71 q	72 r	73 s	74 t	75 u	76 v	77 w
78 x	79 y	7a z	7b {	7c	7d }	7e ~	7f del

0x0a = '\n', 0x20 = ' ', 0x09 = '\t', but note no EOF

## Exercise 2: Using 'a'..'z' as indexes

15/33

Write C code that allows you to treat an array like

```
int freq[26];
```

as if it were indexed by 'a'..'z'

Sample usage

```
for (char c = 'a'; c <= 'z'; c++)
    freq[XXX] = 0;
...
for (char c = 'a'; c <= 'z'; c++)
    printf("%s has freq %d\n", c, freq[XXX]);
```

In other words, replace the xxx by an index calculation

## UTF-8 Character Encoding

16/33

UTF-8 uses a variable-length encoding as follows



#bytes	#bits	Byte 1	Byte 2	Byte 3	Byte 4
1	7	0xxxxxxx	-	-	-
2	11	110xxxxx	10xxxxxx	-	-
3	16	1110xxxx	10xxxxxx	10xxxxxx	-
4	21	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx

The 127 1-byte codes are compatible with ASCII

The 2048 2-byte codes include most Latin-script alphabets

The 65536 3-byte codes include most Asian languages

The 2097152 4-byte codes include symbols and emojis and ...

## ... UTF-8 Character Encoding

17/33

UTF-8 examples

ch	unicode	bits	simple binary	UTF-8 binary
\$	U+0024	7	010 0100	00100100
¢	U+00A2	11	000 1010 0010	11000010 10100010
€	U+20AC	16	0010 0000 1010 1100	11100010 10000010 10101100
☐	U+10348	21	0 0001 0000 0011 0100 1000	11110000 10010000 10001101 10001000

Unicode strings can be manipulated in C (e.g. "안녕하세요")

Like other C strings, they are terminated by a 0 byte (i.e. '\0')

## Exercise 3: Measuring UTF-8 Strings

18/33

Write C functions that count

- the number of bytes in a Unicode string
- the number of "symbols" in a Unicode string

Use the function templates

```
int unicodeNbytes(char *str) { ... }
```


```
int unicodeNsymbols(char *str) { ... }
```

## Numeric Data

19/33

Numeric data comes in two major forms

- integer ... subset (range) of the mathematical integers
- floating point ... subset of the mathematical real numbers



## Integer Constants

20/33

Three ways to write integer constants in C

- 42 ... signed decimal (0..9)
- 0x2A ... unsigned hexadecimal (0..F)
- 052 ... signed octal (0..7)

Variations

- 123U ... unsigned int value (typically 32 bits)
- 123L ... long int value (typically 64 bits)
- 123S ... short int value (typically 16 bits)

Invalid constants lie outside the range for their type, e.g.


- 4294967296, -1U, 6666666S, 078

## Unsigned integers

21/33

The **unsigned int** data type

- commonly 32 bits, storing values in the range  $0 \dots 2^{32}-1$



## ... Unsigned integers

22/33

Value interpreted as binary number

E.g. consider an 8-bit unsigned int

$$01001101 = 2^6 + 2^3 + 2^2 + 2^0 = 64 + 8 + 4 + 1 = 77$$

Addition is bitwise with carry

00000001	00000001	01001101	11111111
+ 00000010	+ 00000011	+ 00001011	+ 00000001
-----	-----	-----	-----
00000011	00000100	01011000	00000000

Most machines will also flag the *overflow* in the fourth example

## Exercise 4: Binary↔decimal Conversion

23/33

Convert these 8-bit binary numbers to hexadecimal:

- 00001001, 00001101, 00101010, 00110011, 11001100

Convert these 8-bit binary numbers to decimal:

- 00001001, 00001101, 00101010, 00110011, 11001100

Convert the following decimal numbers to 8-bit binary:


- 15, 64, 99, 200, 256

## Signed integers

24/33

The `int` data type

- commonly 32 bits, storing values in the range  $-2^{31} \dots 2^{31}-1$



### ... Signed integers

25/33

Several possible representations for negative values

- signed magnitude ... first bit is sign, rest are magnitude
- ones complement ... form  $-N$  by inverting all bits in  $N$
- twos complement ... form  $-N$  by inverting  $N$  and adding 1

In all representations, +ve numbers have 0 in leftmost bit

Examples: representations of (8-bit) -5 (where 5 is 00000101)

- 10000101 ... signed magnitude
- 11111010 ... ones complement
- 11111011 ... twos complement

### ... Signed integers

26/33

*Signed magnitude*: Easy to form  $-X$  from  $X$  ... OR in high-order bit

A problem (using 8-bit ints) ...

- what do these numbers represent? 00000000, 10000000

Two zeroes ... one positive, one negative

Another problem:  $x + -x \neq 0$  (mostly) with simple addition

00000011	3	00101010	42	01111111	127
+ 10000011	-3	+ 10101010	-42	+ 11111111	-127
-----		-----		-----	
10000110	!0	11010100	!0	01111110	!0

To fix requires extra hardware in ALU

### ... Signed integers

27/33

*Ones complement:* Easy to form  $-X$  from  $X$  ... NEG all bits

A problem (using 8-bit ints) ...

- what do these numbers represent? 00000000, 11111111

Two zeroes ... one positive, one negative

At least  $x + -x$  is equal to one of the zeroes with simple addition

00000011    3	00101010    42	01111111
+ 11111100   -3	+ 11010101   -42	+ 10000000
-----	-----	-----
11111111   !0	11111111   !0	11111111   -0

### ... Signed integers

28/33

*Twos complement:* to form  $-X$  from  $X$  ... NEG all bits, then add 1

Now have only one representation for zero (00000000)

- $-0 = \sim 00000000 + 1 = 11111111 + 1 = 00000000$

Only one zero value. Also,  $-(-x) = x$

Even better,  $x + -x = 0$  in all cases with simple addition

00000011    3	00101010    42	01111111
+ 11111101   -3	+ 11010110   -42	+ 10000001
-----	-----	-----
00000000    0	00000000    0	00000000    0

Always produces an "overflow" bit, but can ignore this

## Exercise 5: Binary↔decimal Conversion

29/33

What decimal numbers do these 8-bit twos complement numbers represent:

- 10001001, 10001101, 10101010, 10110011, 11001100

Convert the following decimal numbers to 8-bit binary:

- 15, 64, 99, 127, 128

Show signed magnitude, 1's complement and 2's complement

Demonstrate the addition of  $x + -x$ , where  $x$  is

- 5, 20, 64, 99, 127

## Exercise 6: Integer Powers

30/33

C does not have a power operator (e.g. like  $x * y = x^y$ )

Write a function to compute  $x^y$

```
int raise(int x, int y) { ... }
```

Write a specialised version to compute  $2^y$



```
int powOf2(int y) { ... }
```

---

## Pointers

31/33

Pointers represent memory addresses/locations

- number of bits depends on memory size, but typically 32-bits
- data pointers reference addresses in *data/heap/stack* regions
- function pointers reference addresses in *code* region

Many kinds of pointers, one for each data type

- `sizeof(int *) = sizeof(char *)`  
`= sizeof(double *) = sizeof(struct X *)`

Pointer values must be appropriate for data type, e.g.

- `(char *)` ... can reference any byte address
  - `(int *)` ... must have `addr % 4 == 0`
  - `(double *)` ... must have `addr % 8 == 0`
- 

## ... Pointers

32/33

Can "move" from object to object by *pointer arithmetic*

For any pointer `T *p`, `p++` increases `p` by `sizeof(T)`

Examples (assuming 16-bit pointers):

```
char *p = 0x6060; p++; assert(p == 0x6061)
int *q = 0x6060; q++; assert(q == 0x6064)
double *r = 0x6060; r++; assert(r == 0x6068)
```

A common (efficient) paradigm for scanning a string

```
char *s = "a string";
char *c;
// print a string, char-by-char
for (c = s; *c != '\0'; c++) {
    printf("%c", *c);
}
```

---

## Exercise 7: Sum an array of ints

33/33

Write a function

```
int sumOf(int *a, int n) { ... }
```

to sum the elements of array `a[ ]` containing `n` values.

Implement it two ways:

- using the "standard" approach with an index
  - using a pointer that scans the elements
-